

| Group No | Name | Student Email id | % Contribution |
|----------|-----------------------|------------------|----------------|
| 144 | Abhinav Chandravanshi | 2020FC04860 | Full (100%) |
| 144 | Anurag | 2020FA04028 | Full (100%) |

Databases used for comparison to do performance evaluation:

- **Mongo DB**
MongoDB Compass 1.29.5
MongoDB 5.0.5 Enterprise (64 bit)
Mongo DB Command Line Shell 5.0.5
- **PostgreSQL**
PGAdmin 6.1
PostgreSQL DB 14
PSQL Command Line Shell 14.1

Hardware and OS:

Processor: Intel Core i-3 7100 CPU @ 2.40 GHz
RAM: 6 GB
Storage Type: SSD
Operating System: Windows 10 Home
Architecture Type: 64-bit operating system, x64-based processor

Compass UI

DB Size and Index Size

MongoDB Compass - localhost:27017/restaurant

Connect View Help

Local

4 DBS 2 COLLECTIONS

HOST localhost:27017

CLUSTER Standalone

EDITION MongoDB 5.0.5 Enterprise

Filter your data

- > admin
- > config
- > local

Collections

CREATE COLLECTION

| Collection Name | Documents | Avg. Document Size | Total Document Size | Num. Indexes | Total Index Size | Properties |
|-----------------|-----------|--------------------|---------------------|--------------|------------------|------------|
| train_full | 5,802,400 | 1.7 KB | 9.9 GB | 2 | 151.8 MB | |

Sample documents

MongoDB Compass - localhost:27017/restaurant.train_full

Connect View Collection Help

Local

4 DBS 2 COLLECTIONS

HOST localhost:27017

CLUSTER Standalone

EDITION MongoDB 5.0.5 Enterprise

Filter your data

- > admin
- > config
- > local
- > restaurant

restaurant.train_full

Documents Aggregations Schema Explain Plan Indexes Validation

DOCUMENTS 5.8m TOTAL SIZE 9.9GB AVG. SIZE 1.7KB INDEXES 2 TOTAL SIZE 151.8MB AVG. SIZE 73.9MB

FILTER { field: "value" } OPTIONS FIND RESET ...

ADD DATA VIEW

Displaying documents 1 - 20 of N/A

```
_id:ObjectId("61ca2a00b05fd0001bc3ffc")
customer_Id:"TOMPBT"
gender:"Male"
status_x:Binary('NQ==', 0)
verified_x:Binary('NQ==', 0)
created_at_x:"2018-02-07 19:16:23"
updated_at_x:"2018-02-07 19:16:23"
location_number:
location_type:"work"
latitude_x:-96.44
longitude_x:-67.2
id:20
authentication_id:118616
latitude_y:-0.4075
longitude_y:-71.06
vendor_category_en:"Restaurants"
vendor_category_id:2
delivery_charge:0
serving_distance:8
is_open:true
OpeningTime:"08:00AM-10:45PM"
OpeningTime2:""
preparation_time:19
commission:0
is_acked_delivering:true
discount_percentage:0
status_y:1
verified_y:Binary('NQ==', 0)
rank:1
language:"EN"
vendor_rating:4.5
sunday_from_time1:"08:00:00"
sunday_to_time1:"22:45:00"
```

MONGOSH

Aggregations

MongoDB Compass - localhost:27017/restaurant.train_full

Connect View Collection Help

Local

4 DBS 2 COLLECTIONS

HOST localhost:27017

CLUSTER Standalone

EDITION MongoDB 5.0.5 Enterprise

Filter your data

> admin

> config

> local

> restaurant

train_full

Aggregations

restaurant.train_full

Documents Aggregations Schema Explain Plan Indexes Validation

COLLATION Untitled SAVE SAMPLE MODE AUTO PREVIEW

5802400 Documents in the Collection Preview of Documents in the Collection

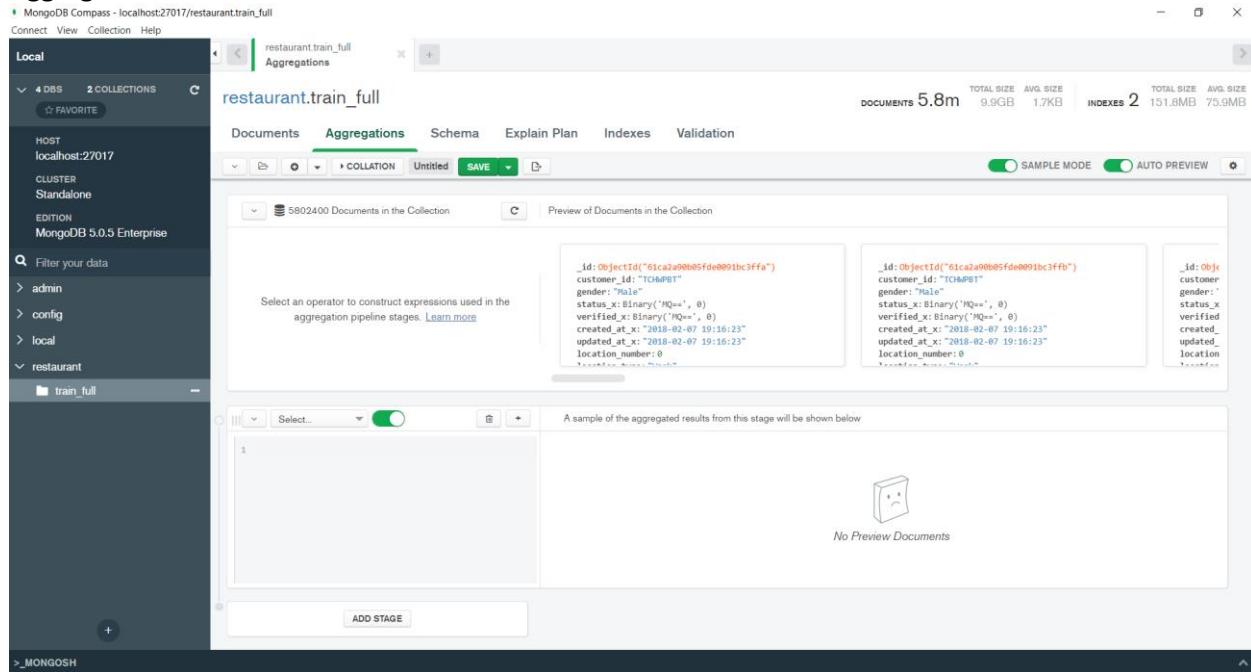
Select an operator to construct expressions used in the aggregation pipeline stages. [Learn more](#)

A sample of the aggregated results from this stage will be shown below

ADD STAGE

No Preview Documents

DOCUMENTS 5.8m TOTAL SIZE 9.9GB AVG. SIZE 1.7KB INDEXES 2 TOTAL SIZE 151.8MB AVG. SIZE 75.9MB



Schema Analysis

MongoDB Compass - localhost:27017/restaurant.train_full

Connect View Collection Help

Local

4 DBS 2 COLLECTIONS

HOST localhost:27017

CLUSTER Standalone

EDITION MongoDB 5.0.5 Enterprise

Filter your data

> admin

> config

> local

> restaurant

train_full

Schema

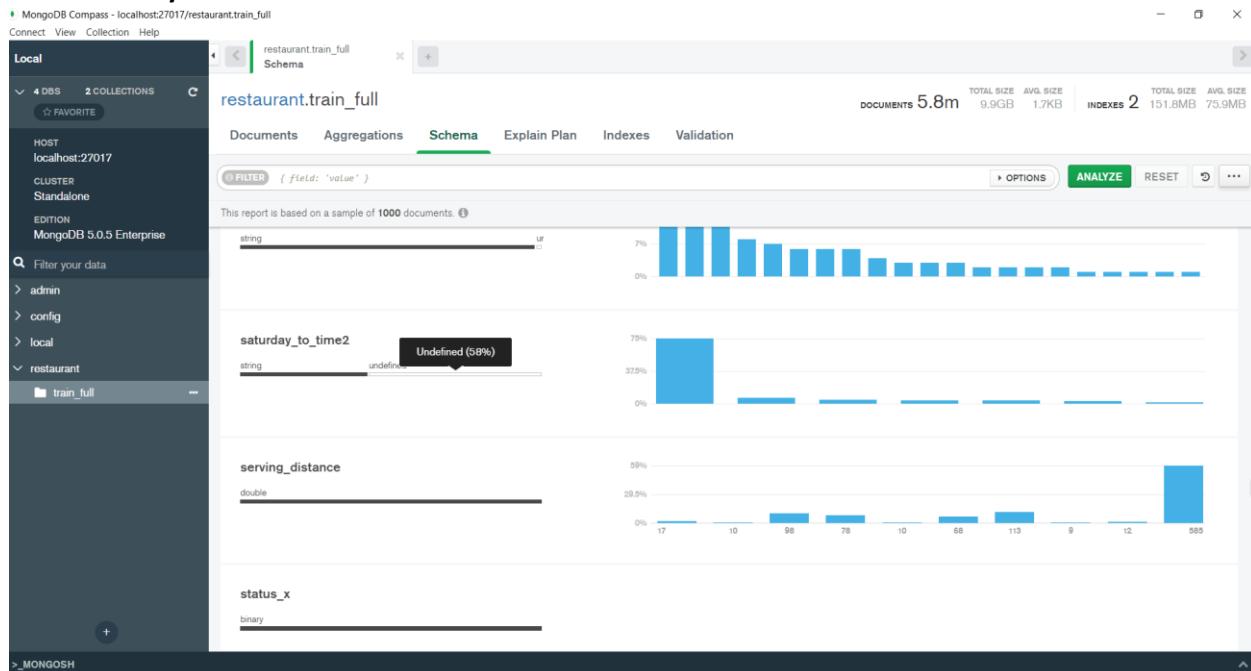
restaurant.train_full

Documents Aggregations Schema Explain Plan Indexes Validation

FILTER { field: 'value' } ANALYZE RESET

This report is based on a sample of 1000 documents.

| Field | Type | Value | Percentage |
|--------|--------|-------------------|------------|
| string | string | uf | 7% |
| string | string | saturday_to_time2 | 37.0% |
| double | double | serving_distance | 59.5% |
| binary | binary | status_x | 0% |



Index

The screenshot shows the MongoDB Compass interface. The left sidebar lists databases (Local, HOST: localhost:27017, CLUSTER: Standalone, EDITION: MongoDB 5.0.5 Enterprise) and collections (admin, config, local, restaurant, train_full). The main area is titled 'restaurant.train_full' and shows the 'Indexes' tab selected. It lists two indexes:

| Name and Definition | Type | Size | Usage | Properties |
|--------------------------|---------|---------|-------------------------|------------|
| CID X LOC_NUM X VENDOR_1 | REGULAR | 78.9 MB | 0 since Mon Dec 27 2021 | UNIQUE |
| _id_ | REGULAR | 72.9 MB | 1 since Mon Dec 27 2021 | UNIQUE |

- **Mongo DB Document Sample Schema**

```
{  
  "created_at_y": "2018-01-30 14:42:04",  
  "created_at_x": "2018-02-07 19:16:23",  
  "gender": "Male",  
  "saturday_from_time2": "10:00:00",  
  "saturday_from_time1": "00:00:00",  
  "_id": {  
    "$oid": "61ca2a90b05fde0091bc3ffa"  
  },  
  "wednesday_from_time2": "08:00:00",  
  "wednesday_from_time1": "00:00:00",  
  "OpeningTime": "11:00AM-11:30PM",  
  "friday_to_time2": "23:59:00",  
  "friday_to_time1": "00:30:00",  
  "target": 0,  
  "sunday_from_time2": "08:00:00",  
  "sunday_from_time1": "00:00:00",  
  "city_id": 1,  
  "delivery_charge": 0,  
  "monday_to_time2": "23:59:00",  
  "monday_to_time1": "00:30:00",  
  "authentication_id": 118597,  
  "primary_tags": "{\"primary_tags\":\"4\"},  
  "CID X LOC_NUM X VENDOR": "TCHWPBT X 0 X 4",  
  "discount_percentage": 0,  
  "thursday_to_time2": "23:59:00",  
  "thursday_to_time1": "00:30:00",  
  "display_orders": 1,  
  "vendor_rating": 4.4,
```

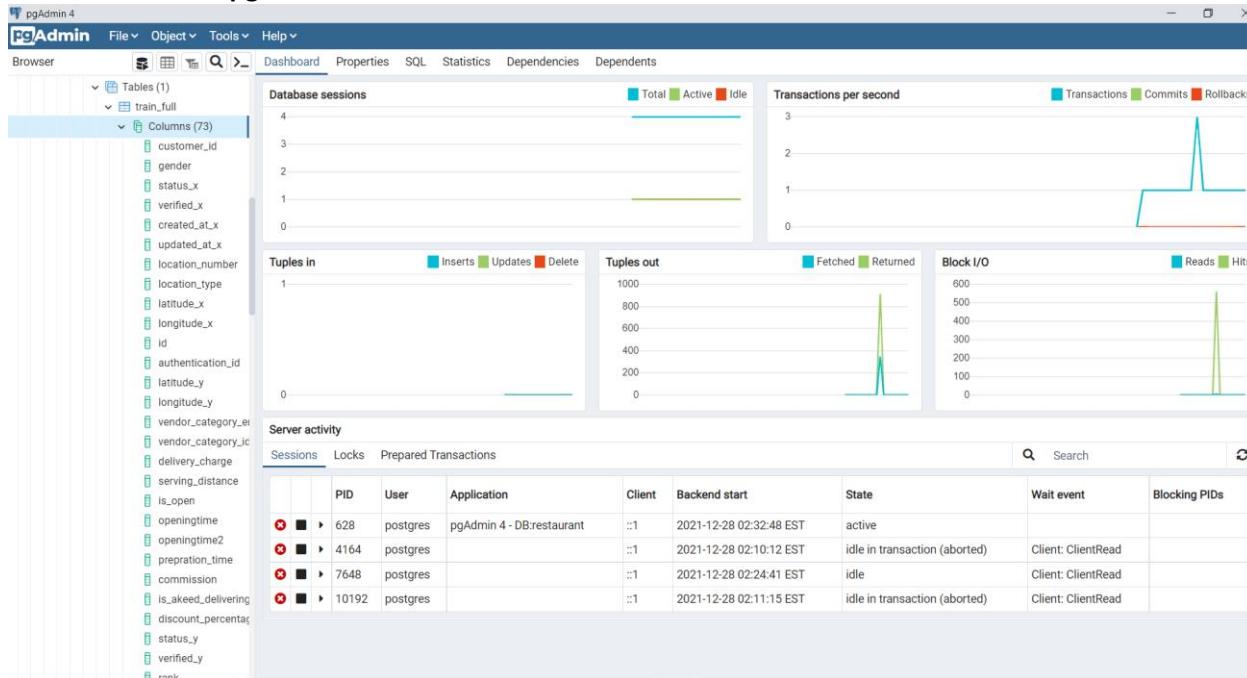
```
"country_id": 1,
"saturday_to_time2": "23:59:00",
"saturday_to_time1": "00:30:00",
"verified_y": {
    "$binary": "MQ==",
    "$type": "0"
},
"verified_x": {
    "$binary": "MQ==",
    "$type": "0"
},
"commission": 0,
"wednesday_to_time2": "23:59:00",
>wednesday_to_time1": "00:30:00",
"is_akeed_delivering": true,
>serving_distance": 6,
"id_obj": 4,
"vendor_tag": "2,4,5,8,91,22,12,24,16,23",
>latitude_y": -0.5884,
>latitude_x": -96.44,
>monday_from_time2": "08:00:00",
>monday_from_time1": "00:00:00",
>open_close_flags": 1,
>device_type": 3,
>location_type": "Work",
>location_number_obj": 0,
>updated_at_y": "2020-04-07 15:12:43",
>updated_at_x": "2018-02-07 19:16:23",
"id": 4,
>is_open": true,
>thursday_from_time2": "08:00:00",
>thursday_from_time1": "00:00:00",
>vendor_category_id": 2,
>vendor_category_en": "Restaurants",
>preparation_time": 15,
>longitude_y": 0.7544,
>longitude_x": -67.2,
>OpeningTime2": "-",
>vendor_tag_name": "Arabic,Breakfast,Burgers,Desserts,Free Delivery,Grills,Lebanese,Salads,Sandwiches,Shawarma",
>status_y": 1,
>status_x": {
    "$binary": "MQ==",
    "$type": "0"
},
>tuesday_to_time2": "23:59:00",
>tuesday_to_time1": "00:30:00",
>customer_id": "TCHWPBT",
>friday_from_time2": "10:00:00",
>friday_from_time1": "00:00:00",
>sunday_to_time2": "23:59:00",
>sunday_to_time1": "00:30:00",
>rank": 11,
>location_number": 0,
```

```

    "tuesday_from_time2": "08:00:00",
    "tuesday_from_time1": "00:00:00",
    "language": "EN",
    "one_click_vendor": true
}

```

Schema loaded in pgAdmin UI



- PostgreSQL Schema

```

CREATE TABLE train_full (
    customer_id VARCHAR,
    gender VARCHAR,
    status_x NUMERIC,
    verified_x NUMERIC,
    created_at_x TIMESTAMP,
    updated_at_x TIMESTAMP,
    location_number NUMERIC,
    location_type VARCHAR,
    latitude_x FLOAT8,
    longitude_x FLOAT8,
    id NUMERIC,
    authentication_id NUMERIC ,
    latitude_y FLOAT8,
    longitude_y FLOAT8,
    vendor_category_en VARCHAR,
    vendor_category_id NUMERIC,
    delivery_charge FLOAT8,
    serving_distance NUMERIC,
    is_open NUMERIC,
    OpeningTime VARCHAR,
    OpeningTime2 VARCHAR,
    preparation_time NUMERIC,
    commission NUMERIC,

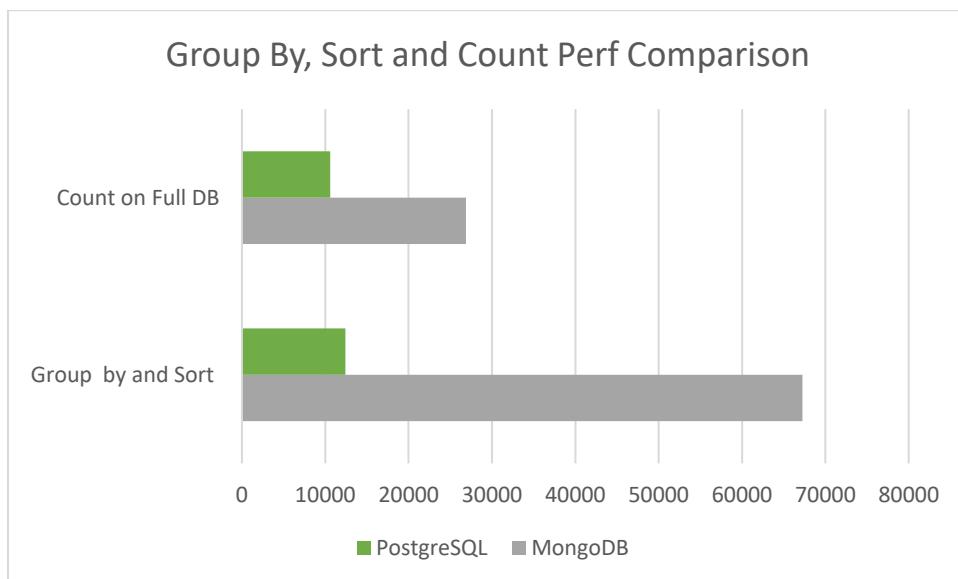
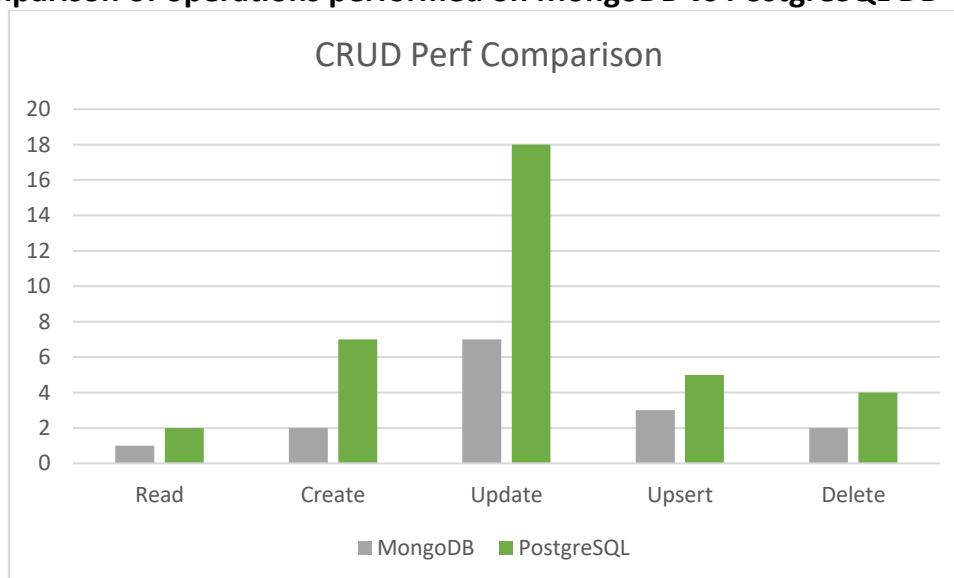
```

```
is_akeed_delivering VARCHAR,
discount_percentage NUMERIC,
status_y NUMERIC,
verified_y NUMERIC ,
rank NUMERIC,
language VARCHAR,
vendor_rating FLOAT8,
sunday_from_time1 TIME,
sunday_to_time1 TIME,
sunday_from_time2 TIME,
sunday_to_time2 TIME,
monday_from_time1 TIME,
monday_to_time1 TIME ,
monday_from_time2 TIME,
monday_to_time2 TIME,
tuesday_from_time1 TIME,
tuesday_to_time1 TIME,
tuesday_from_time2 TIME,
tuesday_to_time2 TIME,
wednesday_from_time1 TIME,
wednesday_to_time1 TIME,
wednesday_from_time2 TIME,
wednesday_to_time2 TIME,
thursday_from_time1 TIME,
thursday_to_time1 TIME,
thursday_from_time2 TIME,
thursday_to_time2 TIME,
friday_from_time1 TIME,
friday_to_time1 TIME,
friday_from_time2 TIME,
friday_to_time2 TIME,
saturday_from_time1 TIME,
saturday_to_time1 TIME,
saturday_from_time2 TIME,
saturday_to_time2 TIME,
primary_tags VARCHAR,
open_close_flags NUMERIC,
vendor_tag VARCHAR,
vendor_tag_name VARCHAR,
one_click_vendor VARCHAR,
country_id NUMERIC,
city_id NUMERIC,
created_at_y TIMESTAMP,
updated_at_y TIMESTAMP,
device_type NUMERIC ,
display_orders NUMERIC,
location_number_obj NUMERIC,
id_obj NUMERIC,
CID_X_LOC_NUM_X_VENDOR VARCHAR PRIMARY KEY,
target NUMERIC);
```

Tabular Summary and Visualization

| DBs | Read (ms) | Create (ms) | Update (ms) | Upsert (ms) | Delete (ms) | Group by and Sort (ms) | Count on Full DB (ms) |
|------------|-----------|-------------|-------------|-------------|-------------|------------------------|-----------------------|
| MongoDB | 1 | 2 | 7 | 3 | 2 | 67237 | 26888 |
| PostgreSQL | 2 | 7 | 18 | 5 | 4 | 12406 | 10613 |

Visual Comparison of operations performed on MongoDB vs PostgreSQL DB



Performance Summary:

- Initial Data load from given source took 85 minutes for mongoDB and it took only 12 minutes for PostgreSQL. This was done using inbuilt native import utility provided by native DBs. So data migration would comparatively take more time for NoSQL DB compared to RDBMS where data model is predefined.
- As we observe from summary table all CRUD operations executed on MongoDB outperform PostgreSQL DB; so MongoDB seems to perform better on CRUD operations. However, being a document based database Groupby, Sort and Counts seem to take more time on MongoDB.

Recommendation:

Hence on basis of CRUD operations, recommendation would be to select MongoDB. Because these operations are performed often in concurrent user environment. Sort and Group are used in less frequent manner and we can fine tune these operations in different ways.

CRUD and other operations on MongoDB

Install dependencies to connect to mongo db (pymongo)

In [1]: ➔ pip install pymongo

```
Requirement already satisfied: pymongo in c:\users\ram\anaconda3\lib\site-packages (4.0.1)
```

```
Note: you may need to restart the kernel to use updated packages.
```

Import required libraries

In [2]: ➔ # importing the required Libraries
import pymongo
import pprint
import json
import warnings
import time
import pprint
warnings.filterwarnings('ignore')

Get connection and db

In [3]: ➔ # connect to the mongoclient
client = pymongo.MongoClient('mongodb://localhost:27017')

get the database
db = client['restaurant']
collection = db['train_full']

Execute a find statement based on unique key and capture execution duration

READ or FIND or SELECT

```
In [7]: ┆ ## Time taken to find the record
#Start time
t0 = time.time()

#execute query and fetch records
record = collection.find( { "CID X LOC_NUM X VENDOR": "TCHWPBT X 0 X 13" } )

#End time
t1 = time.time()

#Elapsed time
total_n = t1-t0

#convert to milisecs
milliseconds = int(round(total_n * 1000))

#print execution time
print('Execution Time(ms): ' + str(milliseconds))

print('Query Result:')
#print the record
for x in record:
    pprint.pprint(x)
```

```
Execution Time(ms): 1
Query Result:
{'CID X LOC_NUM X VENDOR': 'TCHWPBT X 0 X 13',
 'OpeningTime': '08:30AM-10:30PM',
 'OpeningTime2': '-',
 '_id': ObjectId('61ca2a90b05fde0091bc3ffb'),
 'authentication_id': 118608.0,
 'city_id': 1,
 'commission': 0.0,
 'country_id': 1,
 'created_at_x': '2018-02-07 19:16:23',
 'created_at_y': '2018-05-03 12:32:06',
 'customer_id': 'TCHWPBT',
 'delivery_charge': 0.7,
 'device_type': 3,
 'discount_percentage': 0.0,
 'display_orders': 1,
 'friday_from_time1': '00:00:00',
 'friday_from_time2': '08:00:00',
 'friday_to_time1': '01:30:00',
 'friday_to_time2': '23:59:00',
 'gender': 'Male',
 'id': 13,
 'id_obj': 13,
 'is_akeed_delivering': True,
 'is_open': True,
 'language': 'EN',
 'latitude_x': -96.44,
 'latitude_y': -0.4717,
 'location_number': 0,
 'location_number_obj': 0,
```

```
'location_type': 'Work',
'longitude_x': -67.2,
'longitude_y': 0.7446,
'monday_from_time1': '00:00:00',
'monday_from_time2': '08:00:00',
'monday_to_time1': '01:30:00',
'monday_to_time2': '23:59:00',
'one_click_vendor': True,
'open_close_flags': 1.0,
'prepration_time': 14,
'primary_tags': '{"primary_tags":"7"}',
'rank': 11,
'saturday_from_time1': '00:00:00',
'saturday_from_time2': '08:00:00',
'saturday_to_time1': '01:30:00',
'saturday_to_time2': '23:59:00',
'serving_distance': 5.0,
'status_x': b'1',
'status_y': 1.0,
'sunday_from_time1': '00:00:00',
'sunday_from_time2': '08:00:00',
'sunday_to_time1': '01:30:00',
'sunday_to_time2': '23:59:00',
'target': 0,
'thursday_from_time1': '00:00:00',
'thursday_from_time2': '08:00:00',
'thursday_to_time1': '01:30:00',
'thursday_to_time2': '19:30:00',
'tuesday_from_time1': '00:00:00',
'tuesday_from_time2': '08:00:00',
'tuesday_to_time1': '01:30:00',
'tuesday_to_time2': '23:59:00',
'updated_at_x': '2018-02-07 19:16:23',
'updated_at_y': '2020-04-05 20:46:03',
'vendor_category_en': 'Restaurants',
'vendor_category_id': 2.0,
'vendor_rating': 4.7,
'vendor_tag': '4,41,51,34,27,15,24,16,28',
'vendor_tag_name': 'Breakfast,Cakes,Crepes,Italian,Pasta,Pizzas,Salads,Sandwiches,Soups',
'verified_x': b'1',
'verified_y': b'1',
>wednesday_from_time1': '00:00:00',
>wednesday_from_time2': '08:00:00',
>wednesday_to_time1': '01:30:00',
>wednesday_to_time2': '19:30:00'}
```

Explain information on the execution

```
In [8]: ► record = db.train_full.find( { "CID X LOC_NUM X VENDOR": "TCHWPBT X 0 X 13" }  
  
#print the record  
pprint.pprint(record)  
  
{'command': {'$db': 'restaurant',  
             'filter': {'CID X LOC_NUM X VENDOR': 'TCHWPBT X 0 X 13'},  
             'find': 'train_full'},  
 'executionStats': {'allPlansExecution': [],  
                   'executionStages': {'advanced': 1,  
                                       'alreadyHasObj': 0,  
                                       'docsExamined': 1,  
                                       'executionTimeMillisEstimate': 0,  
                                       'inputStage': {'advanced': 1,  
                                         'direction': 'forwar  
d',  
                                         'dupsDropped': 0,  
                                         'dupsTested': 0,  
                                         'executionTimeMillis  
Estimate': 0,  
                                         'indexBounds': {'CID  
X LOC_NUM X VENDOR': ['["TCHWPBT '  
                           'X '  
                           '0 '  
                           'X '  
                           '13", '  
                           '"TCHWPBT '  
                           'X '  
                           '0 '  
                           'X '  
                           '13"]']},  
                                         'indexName': 'CID X  
                                         'LOC_NU  
                                         M X '  
                                         'VENDOR  
                                         _1',  
                                         'indexVersion': 2,  
                                         'isEOF': 1,  
                                         'isMultiKey': False,  
                                         'isPartial': False,  
                                         'isSparse': False,  
                                         'isUnique': True,  
                                         'keyPattern': {'CID  
X LOC_NUM X VENDOR': 1},  
                                         'keysExamined': 1,  
                                         'multiKeyPaths': {'C
```

```

ID X LOC_NUM X VENDOR': []},
                    'nReturned': 1,
                    'needTime': 0,
                    'needYield': 0,
                    'restoreState': 0,
                    'saveState': 0,
                    'seeks': 1,
                    'stage': 'IXSCAN',
                    'works': 2},
                    'isEOF': 1,
                    'nReturned': 1,
                    'needTime': 0,
                    'needYield': 0,
                    'restoreState': 0,
                    'saveState': 0,
                    'stage': 'FETCH',
                    'works': 2},
                    'executionSuccess': True,
                    'executionTimeMillis': 2,
                    'nReturned': 1,
                    'totalDocsExamined': 1,
                    'totalKeysExamined': 1},
        'explainVersion': '1',
        'ok': 1.0,
        'queryPlanner': {'indexFilterSet': False,
                          'maxIndexedAndSolutionsReached': False,
                          'maxIndexedOrSolutionsReached': False,
                          'maxScansToExplodeReached': False,
                          'namespace': 'restaurant.train_full',
                          'parsedQuery': {'CID X LOC_NUM X VENDOR': {'$eq': 'TCHWPB
T X '
                                         '0 X 1
3'}},
                          'rejectedPlans': [],
                          'winningPlan': {'inputStage': {'direction': 'forward',
                                             'indexBounds': {'CID X LOC
_NUM X VENDOR': ['["TCHWPBT '
'X '
'0 '
'X '
'13",
'"TCHWPBT '
'X '
'0 '
'X '
'13"]']}},
                           'indexName': 'CID X LOC_NU
M X '

```

```
                'VENDOR_1',
                'indexVersion': 2,
                'isMultiKey': False,
                'isPartial': False,
                'isSparse': False,
                'isUnique': True,
                'keyPattern': {'CID X LOC_
NUM X VENDOR': 1},
                'multiKeyPaths': {'CID X L
OC_NUM X VENDOR': []},
                'stage': 'IXSCAN'},
                'stage': 'FETCH'}}},
'serverInfo': {'gitVersion': 'd65fd89df3fc039b5c55933c0f71d647a54510ae',
               'host': 'DESKTOP-F81KN39',
               'port': 27017,
               'version': '5.0.5'},
'serverParameters': {'internalDocumentSourceGroupMaxMemoryBytes': 10485760
0,
                     'internalDocumentSourceSetWindowFieldsMaxMemoryByte
s': 104857600,
                     'internalLookupStageIntermediateDocumentMaxSizeByte
s': 104857600,
                     'internalQueryFacetBufferSizeBytes': 104857600,
                     'internalQueryFacetMaxOutputDocSizeBytes': 104857600,
                     'internalQueryMaxAddToSetBytes': 104857600,
                     'internalQueryMaxBlockingSortMemoryUsageBytes': 10485
7600,
                     'internalQueryProhibitBlockingMergeOnMongoS': 0}}
```

CREATE or INSERT

```
In [9]: ┏━ createRec = {'CID X LOC_NUM X VENDOR': 'Bits_Rec1',
  'OpeningTime': '08:30AM-10:30PM',
  'OpeningTime2': '-',
  #'_id': ObjectId('61ca2a90b05fde0091bc3ffb'),
  'authentication_id': 118608.0,
  'city_id': 1,
  'commission': 0.0,
  'country_id': 1,
  'created_at_x': '2018-02-07 19:16:23',
  'created_at_y': '2018-05-03 12:32:06',
  'customer_id': 'bits_cust',
  'delivery_charge': 0.7,
  'device_type': 3,
  'discount_percentage': 0.0,
  'display_orders': 1,
  'friday_from_time1': '00:00:00',
  'friday_from_time2': '08:00:00',
  'friday_to_time1': '01:30:00',
  'friday_to_time2': '23:59:00',
  'gender': 'Male',
  'id': 13,
  'id_obj': 13,
  'is_akeed_delivering': True,
  'is_open': True,
  'language': 'EN',
  'latitude_x': -96.44,
  'latitude_y': -0.4717,
  'location_number': 0,
  'location_number_obj': 0,
  'location_type': 'Work',
  'longitude_x': -67.2,
  'longitude_y': 0.7446,
  'monday_from_time1': '00:00:00',
  'monday_from_time2': '08:00:00',
  'monday_to_time1': '01:30:00',
  'monday_to_time2': '23:59:00',
  'one_click_vendor': True,
  'open_close_flags': 1.0,
  'prepration_time': 14,
  'primary_tags': '{"primary_tags": "7"}',
  'rank': 11,
  'saturday_from_time1': '00:00:00',
  'saturday_from_time2': '08:00:00',
  'saturday_to_time1': '01:30:00',
  'saturday_to_time2': '23:59:00',
  'serving_distance': 5.0,
  'status_x': b'1',
  'status_y': 1.0,
  'sunday_from_time1': '00:00:00',
  'sunday_from_time2': '08:00:00',
  'sunday_to_time1': '01:30:00',
  'sunday_to_time2': '23:59:00',
  'target': 0,
  'thursday_from_time1': '00:00:00',
  'thursday_from_time2': '08:00:00',
  'thursday_to_time1': '01:30:00',
```

```
'thursday_to_time2': '19:30:00',
'tuesday_from_time1': '00:00:00',
'tuesday_from_time2': '08:00:00',
'tuesday_to_time1': '01:30:00',
'tuesday_to_time2': '23:59:00',
'updated_at_x': '2018-02-07 19:16:23',
'updated_at_y': '2020-04-05 20:46:03',
'vendor_category_en': 'Restaurants',
'vendor_category_id': 2.0,
'vendor_rating': 4.7,
'vendor_tag': '4,41,51,34,27,15,24,16,28',
'vendor_tag_name': 'Breakfast,Cakes,Crepes,Italian,Pasta,Pizzas,Salads,Sandw',
'verified_x': b'1',
'verified_y': b'1',
>wednesday_from_time1': '00:00:00',
>wednesday_from_time2': '08:00:00',
>wednesday_to_time1': '01:30:00',
>wednesday_to_time2': '19:30:00'}
```

```
## Time taken to find and print the results
#Start time
t0 = time.time()

collection.insert_one(createRec)

#End time
t1 = time.time()

#Elapsed time
total_n = t1-t0

#convert to milisecs
milliseconds = int(round(total_n * 1000))

#print execution time
print('Execution Time (ms): ' + str(milliseconds))
```

```
Execution Time (ms): 25
```

UPDATE only

```
In [10]: myquery = { 'CID X LOC_NUM X VENDOR': 'Bits_Rec1' }
newvalues = { "$set": { 'customer_id': 'bits_cust11' } }

#Start time
t0 = time.time()

collection.update_one(myquery, newvalues)

#End time
t1 = time.time()

#Elapsed time
total_n = t1-t0

#convert to milisecs
milliseconds = int(round(total_n * 1000))

#print execution time
print('Execution Time(ms) : ' + str(milliseconds))
```

Execution Time(ms) : 7

UPDATE or INSERT if record is not available (UPSERT)

```
In [11]: myquery = { 'CID X LOC_NUM X VENDOR': 'Bits_Rec1' }
newvalues = { "$set": { 'customer_id': 'bits_cust11' } }

## Time taken to find and print the results
#Start time
t0 = time.time()

collection.update_one(myquery, newvalues, upsert=True)

#End time
t1 = time.time()

#Elapsed time
total_n = t1-t0

#convert to milisecs
milliseconds = int(round(total_n * 1000))

#print execution time
print('Execution Time(ms) : ' + str(milliseconds))
```

Execution Time(ms) : 3

DELETE

```
In [12]: ┆ myquery = { 'CID X LOC_NUM X VENDOR': 'Bits_Rec1' }

#Start time
t0 = time.time()

collection.delete_one(myquery)

#End time
t1 = time.time()

#Elapsed time
total_n = t1-t0

#convert to milisecs
milliseconds = int(round(total_n * 1000))

#print execution time
print('Execution Time (ms): ' + str(milliseconds))
```

Execution Time (ms): 4

GROUP BY and SORT

```
In [13]: myquery =[{
    "$group":
    {
        "_id": "$vendor_category_en",
        "avgservdist": { "$avg": "$serving_distance" }
    },
    {
        "$sort": { "avgservdist": 1 }
    }
]

#Start time
t0 = time.time()

agg_result= collection.aggregate(
    myquery)

#End time
t1 = time.time()

#Elapsed time
total_n = t1-t0

#convert to milisecs
milliseconds = int(round(total_n * 1000))

#print execution time
print('Execution Time (ms): ' + str(milliseconds))
print('Query Result:')

for i in agg_result:
    print(i)
```

Execution Time (ms): 67237
 Query Result:
 {'_id': 'Restaurants', 'avgservdist': 11.590909090909092}
 {'_id': 'Sweets & Bakes', 'avgservdist': 12.916666666666666}

COUNT all documents

```
In [14]: #Start time
t0 = time.time()

doc_count = collection.count_documents({})
#End time
t1 = time.time()

#Elapsed time
total_n = t1-t0

#convert to milisecs
milliseconds = int(round(total_n * 1000))

#print execution time
print('Execution Time (ms): ' + str(milliseconds))

print("Total documents:" + str(doc_count))
```

Execution Time (ms): 26888

Total documents:5802400

CRUD and other operations on PostgreSQL DB

Install dependencies

In [1]: ➔ pip install psycopg2

```
Requirement already satisfied: psycopg2 in c:\users\ram\anaconda3\lib\site-packages (2.9.2)
```

```
Note: you may need to restart the kernel to use updated packages.
```

Import libraries

In [2]: ➔ import psycopg2
import time
import pprint

Make DB Connection and get cursor

In [10]: ➔ conn = psycopg2.connect(
host="localhost",
database="restaurant",
user='postgres',
password='admin',
port=5432)
cursor = conn.cursor()

Execute a SELECT statement based on unique key and capture execution duration

READ or FIND or SELECT

```
In [6]: ┏ ## Time taken to find the record
#prepare query
sql_query = "select * from train_full where CID_X_LOC_NUM_X_VENDOR = 'TCHWPBT'

#Start time
t0 = time.time()

#execute query and fetch records
cursor.execute(sql_query)
record = cursor.fetchall()

#End Time
t1 = time.time()

#Elapsed time
total_n = t1-t0

#convert to milisecs
milliseconds = int(round(total_n * 1000))

print('Execution Time (ms): ' + str(milliseconds))

print('Query Result:')
#print the record
for x in record:
    pprint.pprint(x)
```

```
Execution Time (ms): 2
Query Result:
('TCHWPBT',
 'Male',
 Decimal('1'),
 Decimal('1'),
 datetime.datetime(2018, 2, 7, 19, 16, 23),
 datetime.datetime(2018, 2, 7, 19, 16, 23),
 Decimal('0'),
 'Work',
 -96.44,
 -67.2,
 Decimal('13'),
 Decimal('118608.0'),
 -0.4717,
 0.7446,
 'Restaurants',
 Decimal('2.0'),
 0.7,
 Decimal('5.0'),
 Decimal('1.0'),
 '08:30AM-10:30PM',
 '-',
 Decimal('14'),
 Decimal('0.0'),
 'Yes',
 Decimal('0.0'),
 Decimal('1.0'),
 Decimal('1'))
```

```
Decimal('11'),
'EN',
4.7,
datetime.time(0, 0),
datetime.time(1, 30),
datetime.time(8, 0),
datetime.time(23, 59),
'{primary_tags:7}',
Decimal('1.0'),
'4,41,51,34,27,15,24,16,28',
'Breakfast,Cakes,Crepes,Italian,Pasta,Pizzas,Salads,Sandwiches,Soups',
'Y',
Decimal('1.0'),
Decimal('1.0'),
datetime.datetime(2018, 5, 3, 12, 32, 6),
datetime.datetime(2020, 4, 5, 20, 46, 3),
Decimal('3'),
Decimal('1'),
Decimal('0'),
Decimal('13'),
'TCHWPBT X 0 X 13',
Decimal('0'))
```

Explain information on the execution

```
In [7]: ┆ #prepare query
sql_queryexplain = "explain analyze select * from train_full where CID_X_LOC_"

#execute query and fetch records
cursor.execute(sql_queryexplain)
record = cursor.fetchall()

#print the record
pprint.pprint(record)
```

```
[('Index Scan using train_full_pkey on train_full  (cost=0.56..8.57 rows=1
   width=554) (actual time=0.069..0.070 rows=1 loops=1)',),
 (" Index Cond: ((cid_x_loc_num_x_vendor)::text = 'TCHWPBT X 0 X 13'::text"),
 ('Planning Time: 0.264 ms'),
 ('Execution Time: 0.145 ms')]
```

CREATE or INSERT

```
In [11]: ┆ query = """INSERT INTO train_full(customer_id, cid_x_loc_num_x_vendor) VALUES
data = ('Bits_cust','bits_rec2')

#Start time
t0 = time.time()

cursor.execute(query, data)
conn.commit()

#End Time
t1 = time.time()

#Elapsed time
total_n = t1-t0

#convert to milisecs
milliseconds = int(round(total_n * 1000))

print('Execution Time (ms): ' + str(milliseconds))
```

Execution Time (ms): 7

UPDATE only

```
In [12]: ┆ #Start time
t0 = time.time()

cursor.execute("Update train_full set customer_id = ('Bits_custupd') where ci
conn.commit()
#End Time
t1 = time.time()

#Elapsed time
total_n = t1-t0

#convert to milisecs
milliseconds = int(round(total_n * 1000))

print('Execution Time (ms): ' + str(milliseconds))
```

Execution Time (ms): 18

UPDATE or INSERT if record is not available (UPSERT)

```
In [15]: ┆ insert_sql = '''INSERT INTO train_full (customer_id, cid_x_loc_num_x_vendor)
data = ('Bits_cust','bits_rec1')

#Start time
t0 = time.time()

cursor.execute(insert_sql, data)
conn.commit()

#End Time
t1 = time.time()

#Elapsed time
total_n = t1-t0

#convert to milisecs
milliseconds = int(round(total_n * 1000))

print('Execution Time (ms): ' + str(milliseconds))
```

Execution Time (ms): 2

DELETE

```
In [16]: ┆ #Start time
t0 = time.time()

cursor.execute("DELETE FROM train_full where cid_x_loc_num_x_vendor = 'bits_r")
conn.commit()
#End Time
t1 = time.time()

#Elapsed time
total_n = t1-t0

#convert to milisecs
milliseconds = int(round(total_n * 1000))

print('Execution Time (ms): ' + str(milliseconds))
```

Execution Time (ms): 4

GROUP BY and SORT

```
In [17]: ┆ query = """SELECT vendor_category_en, AVG (serving_distance) AS serving_distanc
#Start time
t0 = time.time()

cursor.execute(query)
record = cursor.fetchall()
#End Time
t1 = time.time()

#Elapsed time
total_n = t1-t0

#convert to milisecs
milliseconds = int(round(total_n * 1000))

print('Execution Time (ms): ' + str(milliseconds))

print('Query Result:')
#print the record
for x in record:
    pprint.pprint(x)
```

Execution Time (ms): 12406
 Query Result:
 ('Restaurants', Decimal('11.59090909090909'))
 ('Sweets & Bakes', Decimal('12.916666666666667'))
 (None, None)

Count all records

```
In [18]: ┆ #Start time
t0 = time.time()

cursor.execute("select count(*) from train_full")
#End time
t1 = time.time()

#Elapsed time
total_n = t1-t0

#convert to milisecs
milliseconds = int(round(total_n * 1000))

#print execution time
print('Execution Time (ms): ' + str(milliseconds))

result = cursor.fetchone()
for r1 in result:
    print("Total Records: " + str(r1))
```

```
Execution Time (ms): 10613
Total Records: 5802401
```